

Hadoop - Distributed Computing for the Masses

Casey Stella

November 5, 2012

Table of Contents

Overview

Mesos Project

- ▶ Cluster Dynamic Resource Management
 - ▶ Supports various Frameworks (e.g. Hadoop, MPI, Hypertable, Spark)

Mesos Project

- ▶ Cluster Dynamic Resource Management
 - ▶ Supports various Frameworks (e.g. Hadoop, MPI, Hypertable, Spark)
 - ▶ Multi-resource scheduling (memory and CPU aware)

Mesos Project

- ▶ Cluster Dynamic Resource Management
 - ▶ Supports various Frameworks (e.g. Hadoop, MPI, Hypertable, Spark)
 - ▶ Multi-resource scheduling (memory and CPU aware)
- ▶ Apache Incubator Project (since January 2011)
- ▶ Build new cluster computing frameworks without reinventing low-level facilities for farming out tasks, and have them coexist with existing ones.